

NZ Politician Scorecards — project overview

A data-driven, evidence-linked way to hold New Zealand politicians accountable for how they conduct themselves — not what they believe.

Live demo: <https://nz-politician-scorecards.onrender.com/> (free hosting — the first visit after a quiet period may take ~30s to wake up.)

The problem

Politicians make promises, predictions, and claims with little systematic accountability. Fact-checks are one-off and quickly forgotten; there is no running, public record of *how* a politician conducts themselves over time — whether they answer the question, argue in good faith, keep their promises, or mislead.

What we've built (working prototype)

A platform that scores NZ politicians on **nine conduct attributes** (e.g. Civility, Veracity, Rigour, Specificity, Forthrightness) and — crucially — **links every score to the exact statements behind it**, so anyone can click through and judge the evidence themselves.

- **~85 politicians** scored from **real, scraped public statements**: party press releases, RNZ political reporting, Beehive releases, and **Hansard** (the verbatim parliamentary record).
- **~1,350 individual scored statements**, each kept with its quote, context, and the reasoning for the score.
- Scores are produced by a large language model (Claude) against a fixed, politically-neutral rubric, and measured against **held-out, human-labelled test sets** (e.g. civility rank-correlation ≈ 0.88 on the current set).
- A public website (filter/sort by party, attribute, rarity) and an accompanying card game.

We are honest about the limits

This is a prototype. Accuracy varies by attribute: conduct attributes judgeable from text (civility, specificity) score reliably; attributes that need external verification (veracity, kept-promises) currently reflect *plausibility*, not fact-checked truth. Scores are **automated estimates, not verdicts**, and the evidence is always one click away. We treat this transparency as essential.

Where it needs to go (the principled, scaled version)

To be genuinely useful — ideally **ready for the November 2026 election** — three things need investment:

1. **Scale**: more politicians, more sources (all major outlets), and a deeper time window — years of statements, not months.
2. **Principled methodology**: measure and correct sampling/selection bias; normalise for source mix; surface sample sizes and confidence.
3. **Real accountability for promises & predictions**: a verified, longitudinal **claims ledger** — track every promise/prediction and whether it actually came true, with human review and evidence, rather than an LLM guess.

What we're looking for

- **Funding**, primarily to cover LLM inference at scale (the main running cost).
- **Collaboration** — especially with a newsroom/data desk or researchers — on methodology and editorial standards.

Source, methodology, and evaluation are all open; the site also runs locally in two commands (no API key needed).

Contact: Alexander Telfar — te1farac@gmail.com